

# Mathematical Analysis of a Neural Network for Logistic Regression

Peterson Carara Junior

June, 2025

## Introduction

This work aims to mathematically explain the foundations of training a neural network applied to binary logistic regression, whose objective is to classify asteroids as **hazardous** or **non-hazardous** based on astronomical data from the *Near Earth Objects (NEO)* dataset.

## Problem Description and Dataset

The problem consists of predicting the variable **hazardous** (binary: 0 for non-hazardous, 1 for hazardous) based on variables such as:

- `est_diameter_min`, `est_diameter_max`: Estimated minimum and maximum diameter of the asteroid;
- `relative_velocity`: Relative velocity to the planet;
- `miss_distance`: Approach distance in km;
- `absolute_magnitude`: Absolute magnitude of the celestial body.

After preprocessing, the dataset contains 5 normalized input variables for the neural network.

## Neural Network Architecture

The implemented neural network consists of:

- Fully connected hidden layers with **ReLU** activation function:  $f(x) = \max(0, x)$ ;
- An output layer with **sigmoid** activation:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

which maps the output to the interval  $[0, 1]$ , interpreted as probability.

## Cost Function: Binary Cross-Entropy

The loss function used was binary cross-entropy, given by:

$$\mathcal{L}(y, \hat{y}) = -[y \cdot \log(\hat{y}) + (1 - y) \cdot \log(1 - \hat{y})]$$

Where:

- $y \in \{0, 1\}$ : true label;
- $\hat{y} \in (0, 1)$ : neural network output (probability).

## Optimization: Gradient Descent

During training, the network weights are adjusted to minimize the loss function using the gradient descent algorithm.

### Mathematical Definition

Given a weight vector  $\vec{w}$  and a cost function  $\mathcal{L}(\vec{w})$ , the weight update rule is given by:

$$\vec{w}_{t+1} = \vec{w}_t - \eta \cdot \nabla \mathcal{L}(\vec{w}_t)$$

Where:

- $\eta$ : learning rate (hyperparameter automatically tuned);
- $\nabla \mathcal{L}$ : gradient of the cost function.

The gradient indicates the direction of greatest increase of the loss. Subtracting the gradient means moving in the direction of greatest decrease, seeking the **minimum of the cost function**.

## Hyperparameter Tuning with Random Search

We used the KerasTuner library with Random Search to find the best values for:

- Number of hidden layers (between 1 and 3);
- Number of neurons per layer (between 8 and 64);
- Learning rate  $\eta \in [10^{-4}, 10^{-2}]$ .

## Results

### Classification Report

accuracy:	0.91
precision:	0.92 (Non-Hazardous), 0.73 (Hazardous)
recall:	0.99 (Non-Hazardous), 0.14 (Hazardous)

## Confusion Matrix

$$\begin{bmatrix} 16346 & 93 \\ 1480 & 249 \end{bmatrix}$$

## AUC

Area Under the ROC Curve (AUC) = 0.91

## Graphs




Figure 1: ROC Curve




Figure 2: Accuracy Evolution During Training




Figure 3: Loss Evolution During Training




Figure 4: Gradient vs Dataset Graph

## Final Considerations

The logistic regression implemented through a neural network proved to be efficient in classifying asteroids, despite the dataset imbalance. The use of gradient descent allowed finding a minimum for the cost function, and hyperparameter tuning enhanced the network's ability to learn patterns from the data. From a mathematical perspective, training the network involves vector calculus and partial derivatives to find the minimum of a nonlinear function.